

Remote Method Invocation (RMI)



❖ **Topics**

- Architecture
- Method Parameters
- Simple Example
- Pass by Reference
- Distributed Garbage Collection
- Dynamic Class Loading
- Distributed Polymorphism
- Invocation Traffic

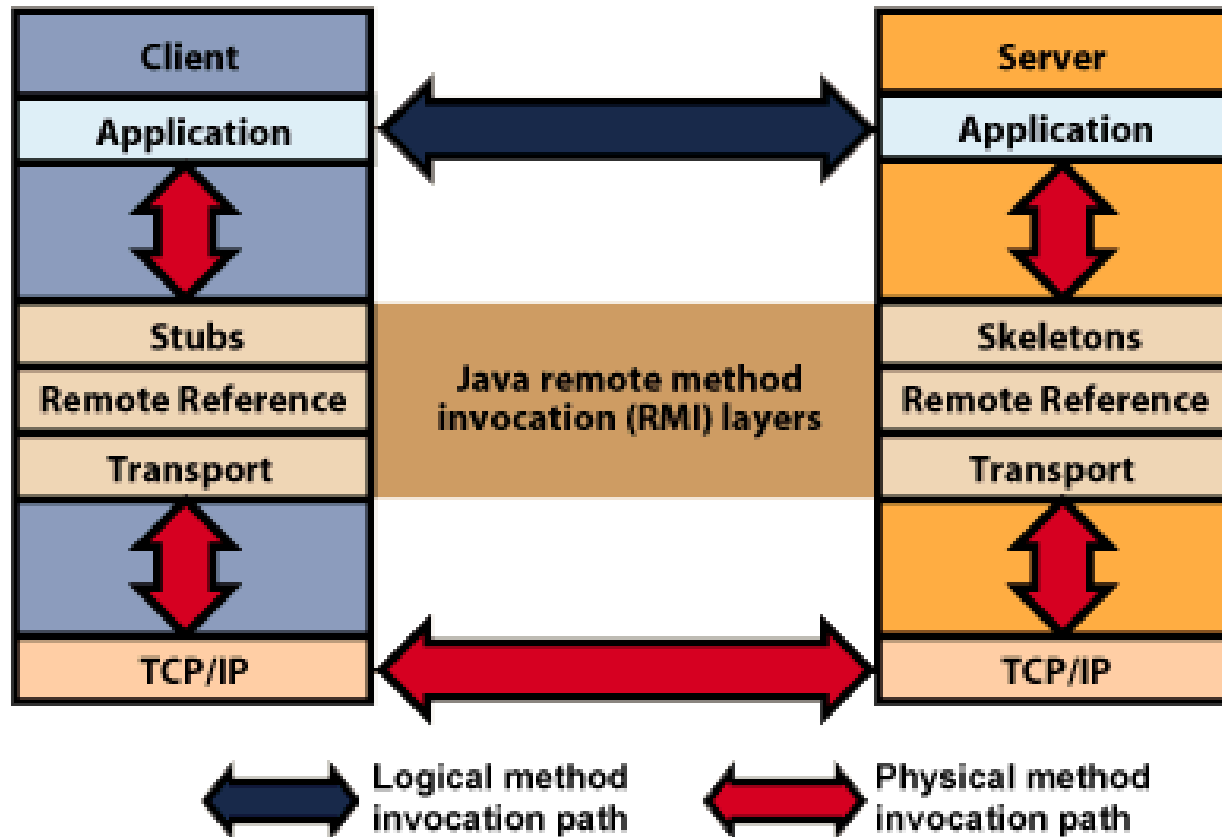
Java Remote Method Invocation

- The Java language's native mechanism for writing distributed objects.
 - **JAVA RMI is Java-centric.**
 - Enables objects to communicate across Java Virtual Machines and physical devices.
- An ***alternative to CORBA, Distributed Component Object Model (DCOM) and SOAP.***
- **RPC**: A procedural invocation from a *process* on one machine to a *process* on another.
 - RMI takes this concept one step further and allows for ***distributed object communication.***
 - RMI allows us to *invoke methods on remote objects.*
 - Can build your networking code as *full objects.*
- Pandora's Box of OO programming opportunities such as *distributed polymorphism.*

RMI Architecture

- Any object whose methods can be invoked from another JVM is called a **Remote Object**.
 - Remote objects are networked objects which expose methods that can be invoked from remote clients.
- Physical location is **not important** (location transparency).
 - Client can invoke methods on an object running on a different VM on the same machine or on a remote machine.
 - Client running in the same address space as a remote object can also invoke its methods.
 - To the remote object, both invocations appear to be the same.
- RMI is a ***rich architecture***:
 - Transparent networking (local/remote transparency).
 - Distributed Garbage Collection.
 - Dynamic Class Loading.
 - Security

RMI Architecture



Interfaces v Implementation

- An interface *defines the exposed information about an object*.
 - The names of its methods, parameters, return types.
- ***The client deals with the interface, not with the remote object.***
 - Interface *masks the implementation* from the client's viewpoint.
 - Clients need only be concerned with the end-result: the methods exposed.
- The ***implementation*** is the core programming logic that an object provides - specific algorithms, logic and data.
- **Why use an interface?**
 - By separating interface from implementation, we can vary an object proprietary logic without changing any client code.
 - E.g. Can plug-in a different algorithm that performs some task more efficiently.


Interfaces v Implementation

- RMI makes *extensive use of interfaces*.
 - All networking code we write is ***applied to an interface***, not an implementation.
 - Impossible to perform a remote invocation directly on an object's implementation.
- *To designate that an object can be accessed remotely, the class must implement the interface `java.rmi.Remote`.*
 - Done by creating our own custom interface extending `java.rmi.Remote`.
- Interface contains definition of each method the remote object exposes.
 - *Each method must also throw a `java.rmi.RemoteException`.*
 - Can be thrown for a variety of reasons – network, registry, class implementation etc...

Stubs and Skeletons

- To mask the fact that you are invoking a remote object, RMI needs to simulate that object locally.
 - Called a ***Stub***.
- *A stub is responsible for accepting method calls locally and ***delegating*** them to their actual object implementations.*
- The remote object in turn accepts calls from a ***skeleton*** that is local to the remote object.
 - ***Skeletons are responsible for receiving calls over the network (from a stub) and delegating them to the remote object implementation.***
- Stubs and skeletons are ***also responsible for marshalling*** – done using Java object serialization.
- ***Note: the use of skeletons is not required from Java2 onwards.***

Stubs and Skeletons



Bootstrapping / RMI Registry

- Before a client and server can start talking, they need some way to connect.
 - ***Acquiring this connection is called bootstrapping.***
- Provided by the ***RMI Registry***.
- To make a remote object accessible, register it with the RMI Registry (yellow pages).
- Registry will then route all incoming requests for that object name to our object.
 - Like a giant hashtable. Just a specific type of naming service.
- When a remote client wants to access an object registered with a particular registry, the client issues a lookup request to the registry.
 - RMI registry is by default ***bound to port 1099*** (can change this if we want).

RMI URLs

- A Java String that is used to located remote objects on another JVM.
- **Conventions:**
 - URL must start with *rmi://*
 - Need to specify the target machine. Defaults to *localhost* if omitted.
 - Append the name of the remote object we want to use:
rmi://www.gmit.ie:1014/myObjectName
- URL is mapped to an object in the RMI Registry using the *java.rmi.Naming* class.
 - Use the ***lookup()*** which takes the URL as a parameter to find an object.
 - Object is initially registered with the RMI Registry using the ***reBind()*** method.
- JNDI is an enterprise extension of the RMI Registry. No need to hard-code URLs – dynamic binding.

The RMIC Compiler

- Stubs and skeletons provide a screen to block networking issues.
 - Both ***must implement*** the remote object's interfaces.
- Use the rmic compiler to generate the stubs and skeletons for us.

>rmic RemoteObjectName

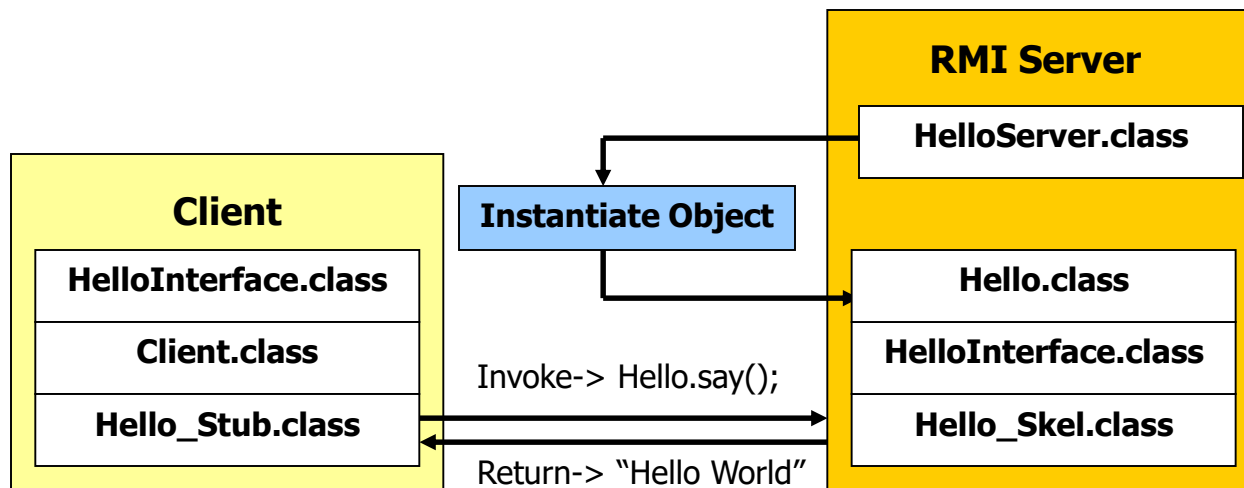
- Can use the *-keepgenerated* option to view the Java source code.
- Note: Since Java 5, it is not necessary to call rmic separately to compile RMI stubs

Object Serialization & Params

- When invoking a method using RMI, all parameters to the remote method are ***passed by value***.
 - i.e. all parameters are copied from one machine to another when the target method is invoked.
- ***Problem:*** If passing an object over the network and that object has references to other objects, how are those references resolved on the target machine?
 - Use object serialization to handle this.
- Serialization is the conversion of a java object to a byte stream representation.
 - Once serialized, it can be sent anywhere – network, disk.
- When you want to use the object again, you must de-serialize it.

Simple Example – Hello World

- We will use RMI to invoke a method called ***say()*** on an object called ***Hello*** and print out the response “Hello World”.
- An overview of the architecture looks like this:



1: Define the Remote Interface

- The Remote interface must have the following properties:
 1. The interface must be ***public***.
 2. The interface must extend the interface `java.rmi.Remote`.
 3. Every method in the interface must declare that it throws `java.rmi.RemoteException`. Other exceptions may also be thrown.

```
import java.rmi.*;
```

```
public interface HelloInterface extends Remote{  
    public String sayHello(String strName) throws RemoteException →  
}  
}
```

2: Implement the Interface

- The implementation of our remote interface is the remote object – the object whose method we wish to invoke.
- The Remote class itself has the following properties:
 1. It must implement a *Remote* interface.
 2. It should extend the `java.rmi.server.UnicastRemoteObject` class. This handles the socket and transport issues.
 - **Question**: Why Unicast?
 - We also can export an object directly in our code using:
`java.rmi.server.UnicastRemoteObject.exportObject()` ;
 - What is the benefit of doing this?
 3. It **can** have methods that are ***not in its Remote interface***. These can only be invoked locally – why?

2: Implement the Interface

```
import java.rmi.*;
import java.rmi.server.*;
public class HelloWorld extends UnicastRemoteObject →
    implements HelloInterface{
    private String strMessage;

    public HelloWorld(String msg) throws RemoteException{
        strMessage = msg;
    }

    public String sayHello(String strName) throws →
RemoteException{
        return strMessage + " " + strName;
    }
}
```


3: Create Stubs and Skeletons

- After compiling the remote object, we can create the stubs and skeletons using the rmic compiler:

```
c:>rmic HelloWorld
```

- This generates the following files:
 - HelloWorld_Skel.class
 - HelloWorld_Stub.class
- Recall that we can use the *-keepgenerated* switch to generate the .java files for our perusal.

4: Instantiate the Remote Object

- So far we have defined and implemented a remote interface.
- But a client can only invoke a method on an object that exists.
 - ***Our object has not been instantiated yet.***
- Also, we need to register our object with the RMI Registry so a client can find it.
- Therefore we use a “server” or factory class to create our remote object for us – recall: an object is an instance of a class.

4: Instantiate the Remote Object

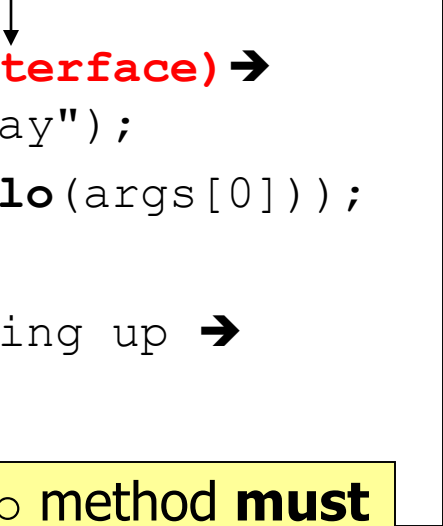
```
import java.rmi.*;
import java.rmi.server.*;

class HelloServer{
    public static void main(String args[]){
        try{
            Naming.rebind("HelloWorld", new →
                HelloWorld("Howday"));
            System.out.println("The HelloWorld server →
                is ready...");
        }catch(Exception e){
            System.out.println("Error creating →
                HelloWorld object.");
        }
    }
}
```

5: Create the Client

```
import java.rmi.*;
import java.rmi.server.*;

public class Client{
    public static void main(String args[]){
        try{
            HelloInterface hello = (HelloInterface) →
Naming.lookup("//localhost/Howday");
            System.out.println(hello.sayHello(args[0]));
        }catch (Exception e){
            System.out.println("Error: looking up →
            object in RMI registry." + e);
        }
    }
}
```



The result of the `Naming.lookup` method **must be cast** to the type of the remote interface

6: Start the RMI Registry

- Before starting the Server, we need to start the Object Registry, and leave it running in the background:

rmiregistry

- It takes a second or so for the Object Registry to start running and to start listening on its socket (what port number?).
- The server should then be started with the java command.
- The client can then be started and make its remote method invocation.

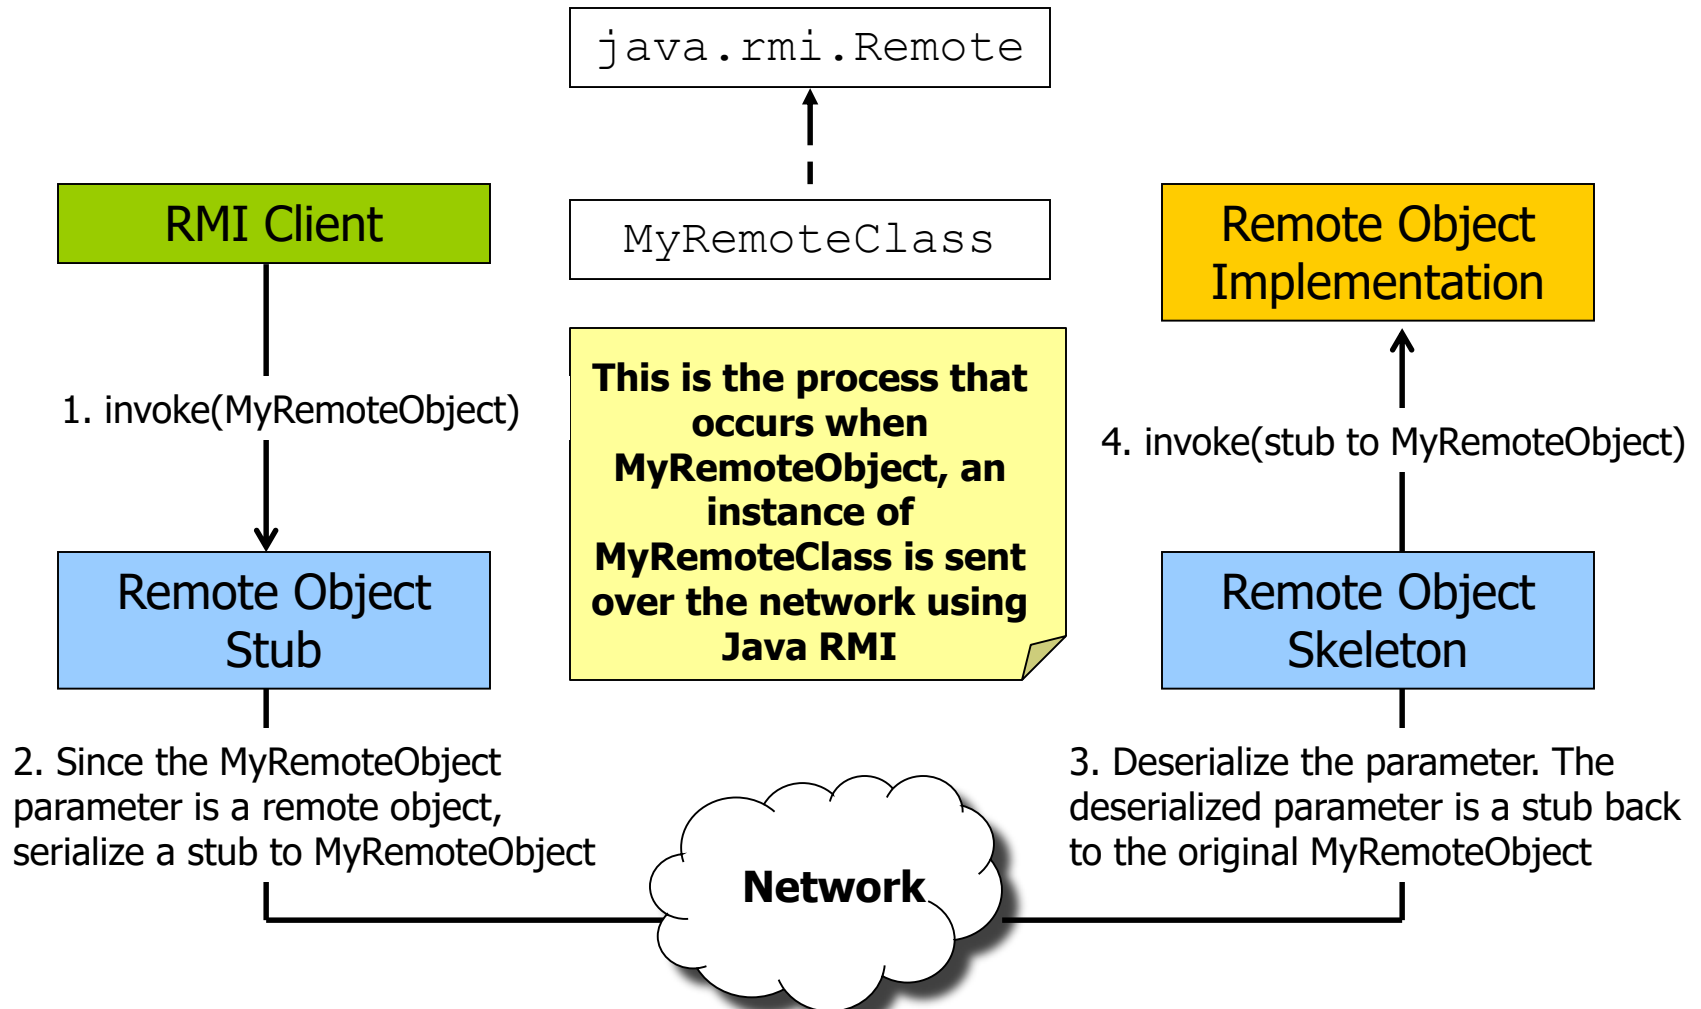
Pass by Reference

- So far we have looked at passing parameters by value.
- An entire graph of objects has to be serialized and sent across the network.
 - What happens if graph is very large?
 - What if we have lots of state information to send?
 - Performance degradation!
- RMI allows us to get around this by simulating a pass by reference – a copy of an object parameter is not sent.
 - The remote method works with the original object on the client.
- ***To do this, the object parameter must itself be a remote object.***
- When the remote host performs an action on the object parameter, this action occurs on the local host.
 - ***i.e. The operation is itself a remote method invocation.***

Pass by Reference

- What is actually passed to the remote object?
 - The stub for our local object (which itself is remote with reference to the remote object).
- In Java we can have references to many objects.
- RMI is an extension of this allowing us references to remote objects.
- `java.rmi.RemoteStub` objects are the manifestation of this.
- They are serializable and can be passed over the network – class loading.
- In summary, Java RMI simulates a pass by reference by passing a serialized stub, rather than passing a serialized object.

Pass by Reference



Distributed Garbage Collection

- RMI uses a reference-counting garbage collection algorithm similar to Modula-3's Network Objects.
 - RMI runtime keeps track of ***all live references*** within each JVM.
- Described by ***java.rmi.dgc.DGC*** interface methods: ***dirty()*** and ***clean()***.
 - ***Dirty call*** is made when a remote reference is unmarshaled in a client (the client is indicated by its VMID). ***Reference count*** is incremented
 - Corresponding ***clean call*** is made when no more references to the remote reference exist in the client.
 - ***As live references are found to be unreferenced in the local VM, the count is decremented.***
- A remote object not referenced by any client is called a ***weak reference***.
 - Allows the JVM's garbage collector to discard the object if no other local references to the object exist.

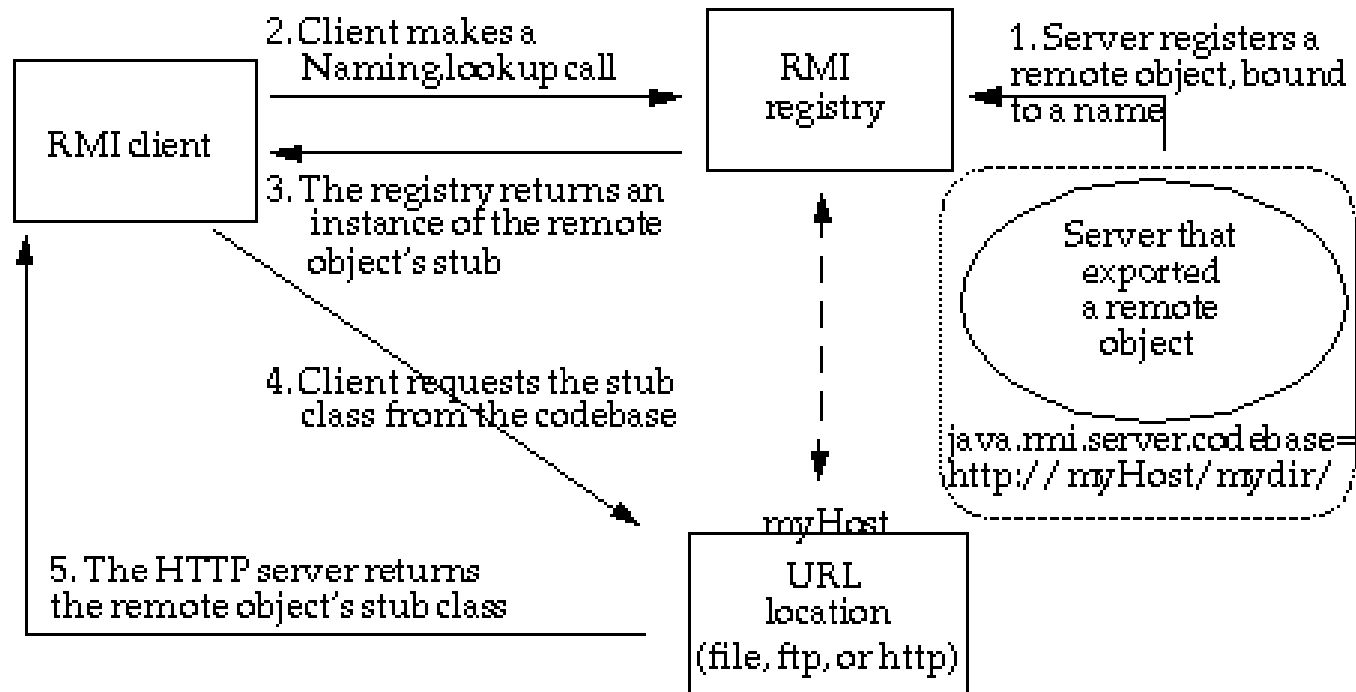
Distributed Garbage Collection

- As long as a local reference to a remote object exists, it cannot be garbage-collected.
 - Can be passed in remote calls or returned to clients.
- Passing a remote object ***adds the identifier*** for the VM to which it was passed to the referenced set.
 - A remote object needing unreferenced notification must implement the ***java.rmi.server.Unreferenced*** interface.
 - ***unreferenced()*** method invoked when references no longer exist.
 - ***finalized()*** (if implemented) will be invoked when the GC is about to reclaim an objects memory space.
- **Note:** References cannot guarantee ***referential integrity***. possible that a remote reference may in fact not refer to an existing object.
 - Transport might think that application has crashed.

Dynamic Class Loading

- RMI allows parameters, return values and exceptions passed in RMI calls to be any object that is **serializable**.
 - Annotates the serialized call stream with the **appropriate location information** so that the class definition files can be loaded at the receiver.
- **Require class definitions** for unmarshalled parameters/return types.
 - Unmarshalling process first attempts to resolve classes by name in its **local CLASSPATH**.
 - Can also **dynamically load class definitions** over a network (including remote stub classes).
- Done using special subclasses of **ObjectOutputStream** and **ObjectInputStream**. (**RMIClassLoader**)
 - Subclasses override the **annotateClass()** of **ObjectOutputStream** and **resolveClass()** of **ObjectInputStream**.
 - Communicates locations for class definitions of *class descriptors in the stream*.

Dynamic Class Loading



Dynamic Class Loading

- Normally dynamically load classes over **HTTP or FTP**.
 - ***java.rmi.server.RMIClassLoader*** performs this function.
- Controlled by **properties** specified when each JVM is run:
java -D<propertyName>=<propertyValue> + className
 - Use ***java.rmi.server.codebase*** property to specify a URL (same as applets).
 - Points to a HTTP/FTP service that supplies classes for objects ***sent from this JVM.***
 - This URL is sent in serialized byte stream to client along with object.
- **Note:** RMI does ***not*** send class files along with serialized objects.
 - If remote JVM needs to load a class file for an object, it looks for the embedded URL and contacts server at that location for the class.
- If ***java.rmi.server.useCodebaseOnly*** property is set to true, JRE will load classes only from location specified by the ***CLASSPATH*** or a URL specified by ***java.rmi.server.codebase*** property.
 - This is an easy way of imposing ***additional security*** to the behaviour of ***RMIClassLoader***.

Class Loading Configurations

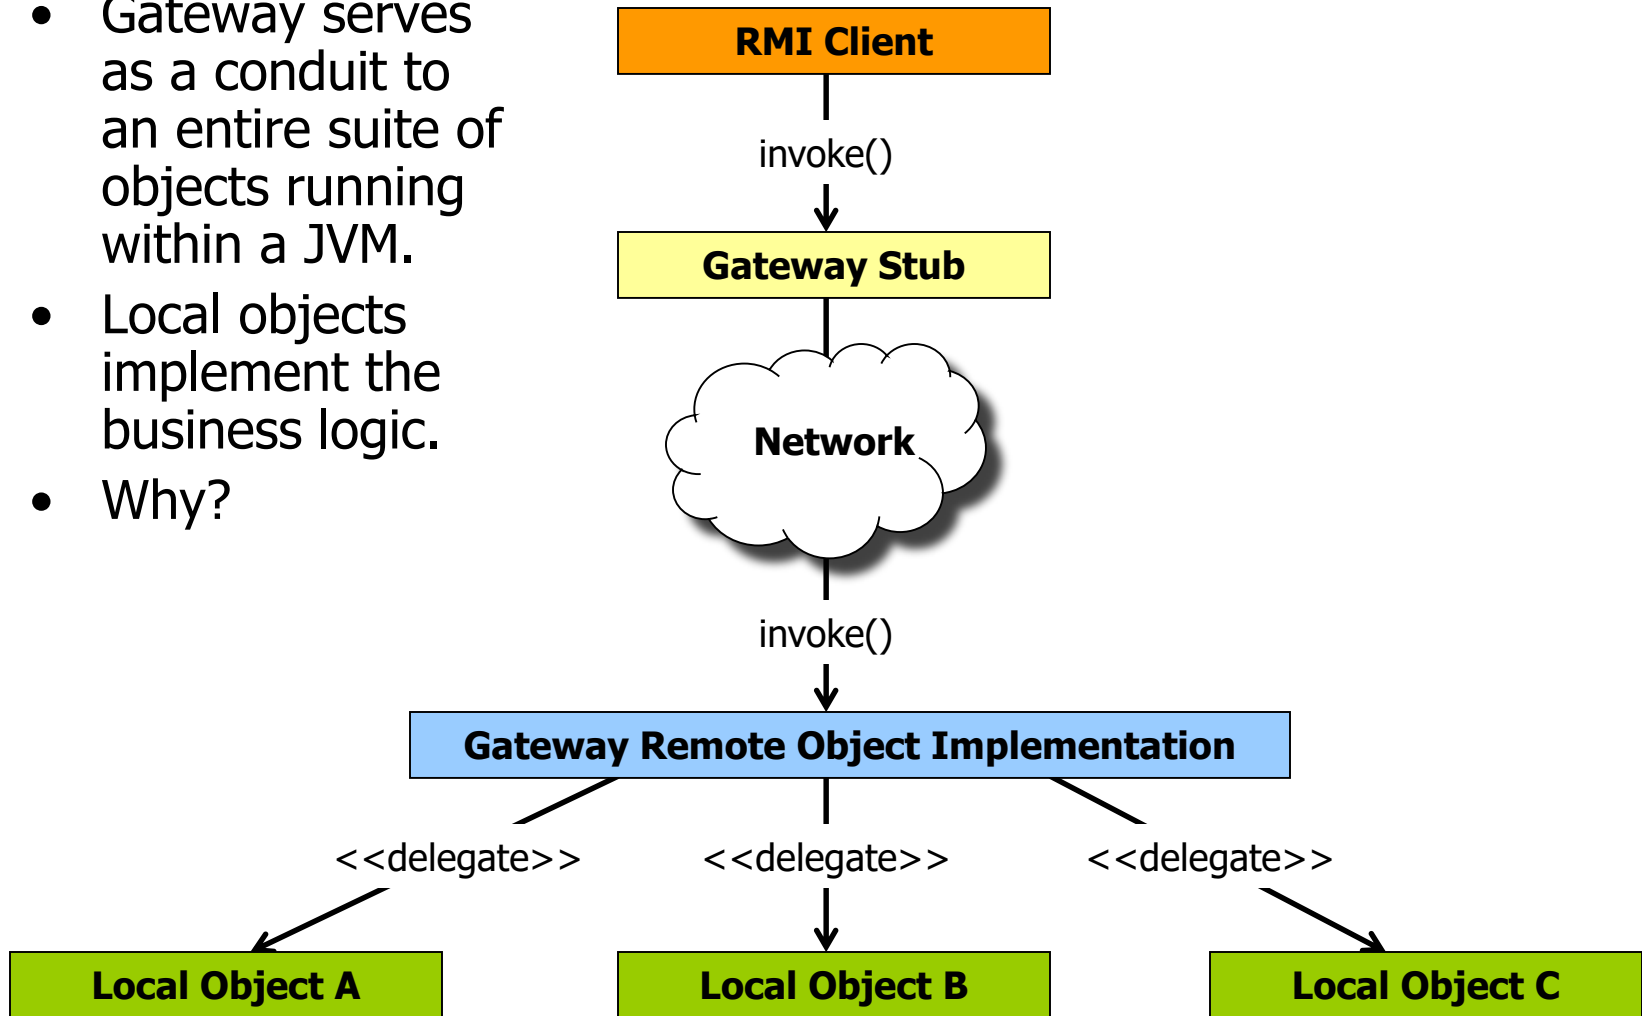
- **Closed**: All classes used by clients and server must be located on JVM and referenced by the *CLASSPATH* environment variable. No dynamic class loading is supported.
- **Server based**: Client is loaded from the server's CODEBASE along with all supporting classes. Similar to the way applets are loaded.
- **Client dynamic**: Primary classes loaded by referencing the *CLASSPATH* environment variable of the JVM for the *client*. Supporting classes loaded by *java.rmi.server.RMIClassLoader* via HTTP/FTP from *server specified* location.
- **Server-dynamic**: Primary classes loaded by referencing the *CLASSPATH* environment variable of the JVM for the *server*. Supporting classes loaded by *java.rmi.server.RMIClassLoader* via HTTP/FTP from client specified server.

Class Loading Configurations

- **Bootstrap client:** All client code is loaded via HTTP/FTP server across the network. Only code residing on the client machine is a small bootstrap loader.
- **Bootstrap server.** All of the server code is loaded via HTTP/FTP server located on the network. Only code residing on the server machine is a small bootstrap loader.

Using Gateway to Mask RMI

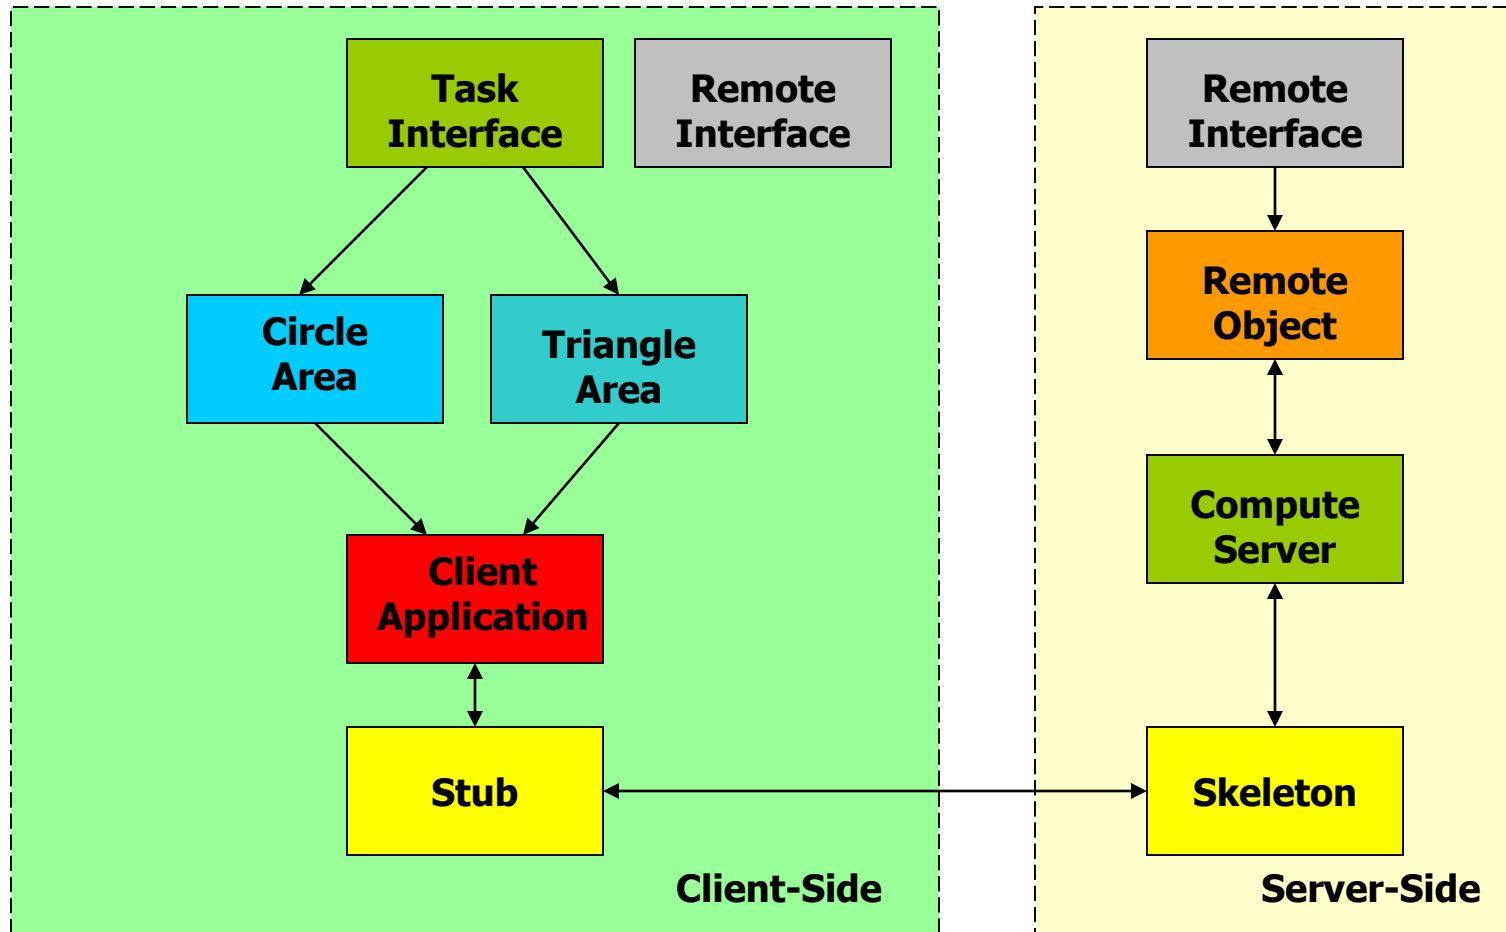
- Gateway serves as a conduit to an entire suite of objects running within a JVM.
- Local objects implement the business logic.
- Why?



Distributed Polymorphism

- ***The ability to recognise at runtime, the actual implementation type for a particular interface.***
- We will use an example of a remote object that calculates arbitrary tasks:
 - A client application sends a task object to a compute server.
 - The compute server runs the task and returns a result.
 - RMI loads the task code dynamically on server.
- This example shows polymorphism on the server.
 - Will also work on the client.
- *Server returns a particular interface implementation.*

Distributed Polymorphism



The Task and Remote Interfaces

```
package gmit;
import java.io.Serializable;

public interface Task extends Serializable{
    public Object run();
}
```

```
package gmit;
import java.rmi.Remote;
import java.rmi.RemoteException;


public interface Compute extends Remote{
    public Object runTask(Task t) throws RemoteException;
}
```

The Remote Object & Server

```
package gmit;
import java.rmi.*;
import java.rmi.server.*;
public class ComputeServer extends UnicastRemoteObject implements Compute{
    public ComputeServer() throws RemoteException{

public Object runTask(Task t) throws RemoteException{
    return t.run();
}

    public static void main(String args[]){
        System.setSecurityManager(new RMISecurityManager());
        try{
            ComputeServer cs = new ComputeServer();
            Naming.rebind("Computer", cs);
        }catch(Exception e){}
    }
}
```



We haven't even written the code that the server is going to execute yet!

The CircleArea Task

```
package gmit;

public class CircleArea implements Task{

    private double radius;

    public Object run(){
        Object value = (Object) new Double(radius * radius * Math.PI);

        return (value);
    }

    public void setRadius(double rad){
        radius = rad;
    }
}
```

The TriangleArea Task

```
package gmit;
```

```
public class TriangleArea implements Task{
```

```
    private double base;
```

```
    private double height;
```

```
    public Object run(){
```

```
        Object value = (Object) new Double((base/2)*height);
```

```
        return (value);
```

```
    }
```

```
    public void setBase(double b){
```

```
        base = b;
```

```
    }
```

```
    public void setHeight(double h){
```

```
        height = h;
```

```
    }
```

```
}
```

The Client Application

```
package gmit;

import java.rmi.*;

public class PolyClient{
    public static void main(String args[]){
        try{
            Compute comp = (Compute) Naming.lookup("///Computer");

            CircleArea circle = new CircleArea();
            circle.setRadius(1234);

            TriangleArea triangle = new TriangleArea();
            triangle.setBase(1234);
            triangle.setHeight(1234);

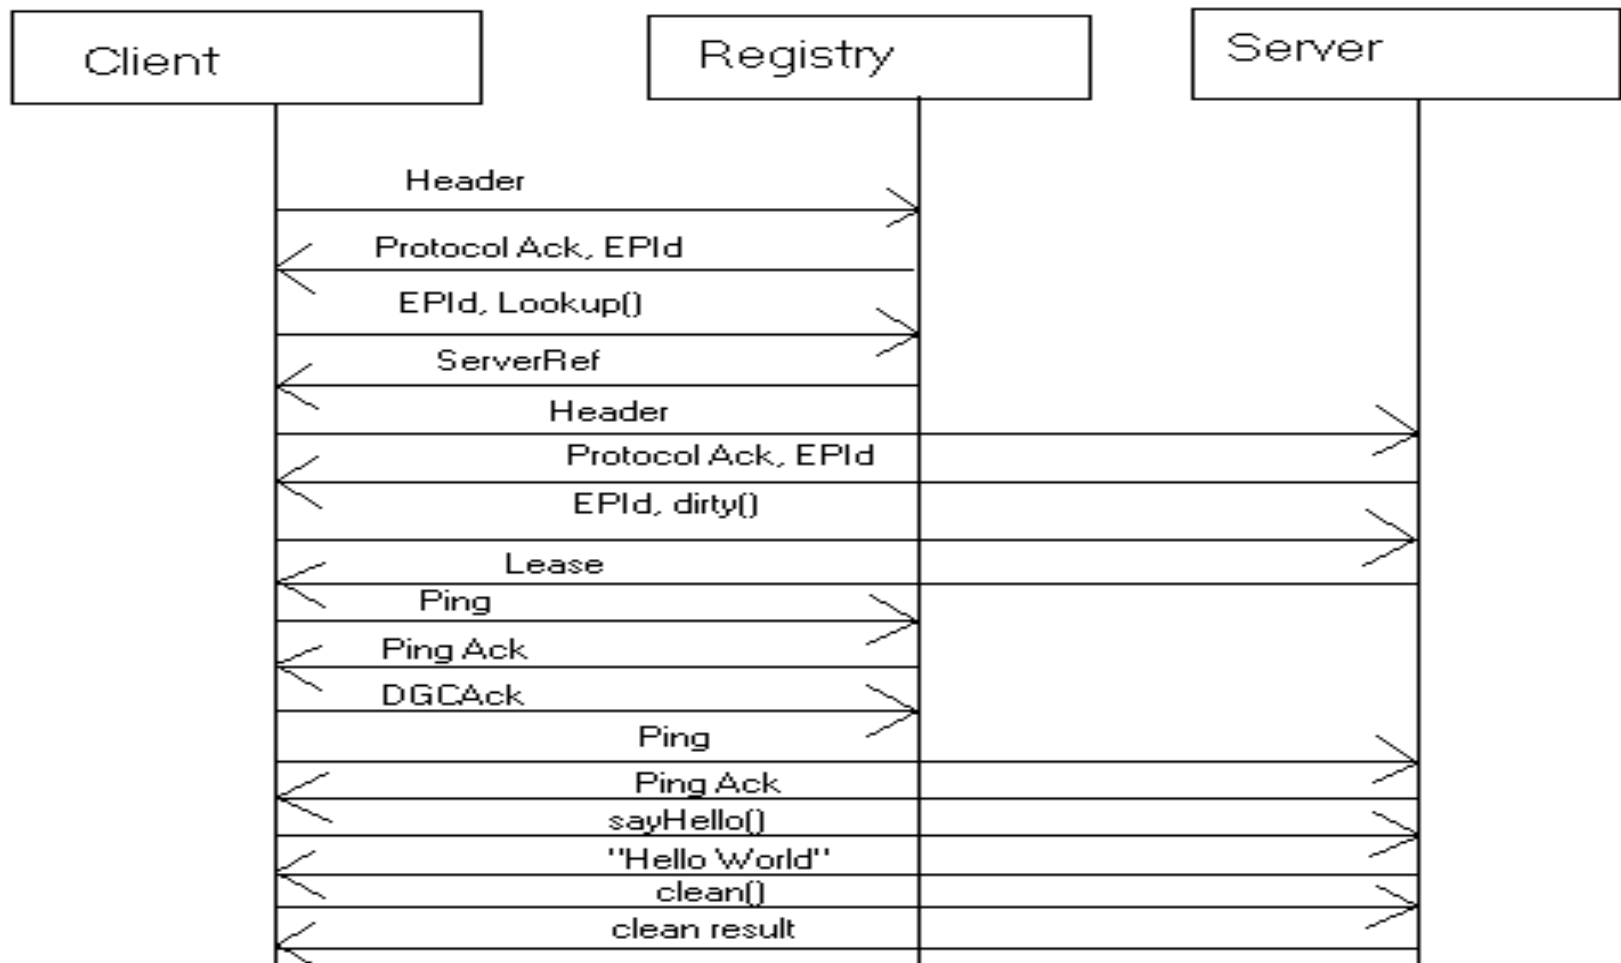
            String circleResult = ((Double) comp.runTask(circle)).toString();
            String triangleResult = ((Double) comp.runTask(triangle)).toString();
```

The Client Application

```
        System.out.println("Area of Circle:" + circleResult);
        System.out.println("Area of Triangle:" + triangleResult);

    }catch(NotBoundException nbe){
        System.out.println("NotBoundException Error: " + nbe.getMessage());
        System.exit(0);
    }catch(RemoteException re){
        System.out.println("RemoteException Error: " + re.getMessage());
        System.exit(0);
    }catch(Exception e){
        System.out.println("General Exception Error: " + e);
        System.exit(0);
    }
}
}
```


Remote Call Message Sequence



Invocation Data Traffic

	Client to server & registry	Server & Registry to Client	Total
Registry Lookup	55(6%)	276(42%)	331(20%)
Invocation Data	41(4%)	37(6%)	78(5%)
DGC Data	831(85%)	305(46%)	1136(69%)
Protocol Overhead	52(5%)	40(6%)	92(6%)
Total	979(100)%	658(100%)	1637(100%)

- From Campadello et al, *"Performance Enhancing Proxies for Java2 RMI over Slow Wireless Links"*, *Proceedings of the Second International Conference on the Practical Application on Java*, 2000.